

# Linux Tracing Tools

Perf and BCC (eBPF)

Ravi Bangoria

ravi.bangoria@linux.ibm.com

Linux Technology Center

India Systems Development Lab, IBM

Sandipan Das

sandipan@linux.ibm.com

Where do you need **Tracing**?


# Perf

## Why Perf?

- Low overhead
- Easy to use
- Production-safe
- Feature rich
- No service daemons are needed
- Doesn't stop the workload
- No need to recompile the workload

# Agenda

- Performance Monitoring
  - Counting events
  - Profiling (Sampling)
  - Callgraph
  - Drill down to instruction level
  - Perf script
  - Real time profiling
- Static Tracing
- Dynamic Tracing (Kprobe/Uprobe)
- Hardware Breakpoints
- BCC (eBPF)



# Counting events

Sub command

Optional

```
# perf stat -e cycles -- <workload>
```

Perf tool

Event

The diagram shows the command `# perf stat -e cycles -- <workload>` with several annotations. A magenta label 'Sub command' has a dashed arrow pointing down to the word 'stat'. A black label 'Optional' has a dashed arrow pointing down to the double hyphen '--'. A blue label 'Perf tool' has a dashed arrow pointing up to the word 'perf'. A green label 'Event' has a dashed arrow pointing up to the word 'cycles'. The word 'cycles' is green, while 'stat' and '<workload>' are magenta.

# Profiling (Sampling)

# perf record <workload>

# perf report

- Record => perf.data => Report
- Default event is cycles. Use -e for other events
- Default frequency is 4k samples/second
- Record with -a for systemwide, -p for pid specific, -C for CPU specific
- Supports cross arch record / report

# Symbol table / Debuginfo

# readelf -SW <binary>

- Map instruction pointer to symbol name
- Install debuginfo package if distro provided binary is stripped
- Does this mean I've to install debuginfo packages on production system? -- NO.

# Callgraph

# perf record -g ...

# perf report --no-children

```
Samples: 17K of event 'cycles:ppp', Event count (approx.): 5462943000
```

	Overhead	Command	Shared Object	Symbol
+	15.53%	swapper	[kernel.vmlinux]	[k] snooze_loop
+	4.02%	swapper	[unknown]	[H] 0xc0000000000eebbc
-	3.58%	swapper	[kernel.vmlinux]	[k] __lock_acquire
-		__lock_acquire		
+	1.90%	0		
-	1.45%	lock_acquire		
-	0.40%	_raw_spin_lock_irqsave		
+	0.19%	hrtimer_get_next_event		
+	0.15%	hrtimer_next_event_without		
-	0.29%	ktime_get		
+	0.15%	tick_nohz_irq_exit		
-	0.12%	tick_irq_enter		
		irq_enter		
-		timer_interrupt		
+	0.10%	plpar_hcall_norets		
+	0.27%	_raw_spin_lock		
+	0.23%	tick_nohz_next_event		
+	0.15%	timekeeping_max_deferment		



# Drilldown at instruction level

## # perf annotate

- With / Without Source
- Interactive
- Dependencies of instructions
- Read my blog:

<https://www.ibm.com/developerworks/library/l-analyzing-performance-perf-annotate-trs/index.html>

```
snooze_loop /lib/modules/4.18.0-rc4+/build/vmlinux
Percent      asm (CURRENT_THREAD_INFO(%0,1) : "=r" (val));
b0: → rldicr r10,r1,0,49
test_bit():
 * @nr: bit number to test
 * @addr: Address to start counting from
 */
static inline int test_bit(int nr, const volatile unsi
{
    return 1UL & (addr[BIT_WORD(nr)] >> (nr & (BIT
    ld      r9,128(r10)
snooze_loop():
    while (!need_resched()) {
        andi. r9,r9,4
        ↓ bne 100      HMT_low();
        mr     r1,r1      HMT_very_low();
        mr     r31,r31
        if (likely(snooze_timeout_en) && get_t
        addis  r9,r2,4
        lbz   r9,8896(r9)
        cmpwi cr7,r9,0
        31.22  beq    cr7,b0
get_tb():
        12.21  mftb   r9
```

# Perf script

# perf script

or

# perf script record/report <script> -- <workload>

- Read perf.data and display trace output
- Python and Perl support
- Use --list option for available scripts
- You can write your own script
- Use -g to generate new template from perf.data

# Real Time Profiling

# perf top

- Generate and display profile in real time
- “Top” like but more detailed
- Supports call-graphs, annotate etc. feature in real time
- Interactive

# Static Tracing (Tracepoints)

- Profiling takes samples. Tracing records every event.
- Tracepoints are added by developer at important places in the code.
- Kernel tracepoints are already supported.
- Userspace tracepoints(USDT) are partially supported. We are working on providing full support.
- When not being traced they are just “nop”.
- Each tracepoints has their own list of arguments and output format.

# Dynamic Tracing (Kprobes)

```
# perf probe <probe_location>
```

```
# perf record -e probe:...
```

- Create dynamic tracepoints in kernel
- Probe with arguments, variables
- Return probes
- Probe inside kernel module
- Probe location can be a function, file:lineno, function+offset ...
- Need debuginfo for some of the features

# Dynamic Tracing (Uprobes)

```
# perf probe -x binary <probe_location>
```

```
# perf record -e probe:...
```

- Create dynamic tracepoints in userspace application
- Probe with arguments, variables
- Return probes
- Probe location can be a function, file:lineno, function+offset ...
- Need debuginfo for some of the features

# Hardware Breakpoints

# perf record -e mem:0xc0000000011ea98c ...

- Types of hw-breakpoints:
  - Data breakpoints (Watchpoints)
  - Instruction breakpoints
- Watchpoint: Who is changing a particular memory location?
- Watchpoints useful for debugging memory corruption problems.
- Instruction breakpoints is same as kprobes but supported by hw.
- Very limited hw registers.


# Linux Tracing with eBPF



# eBPF

## What is eBPF?

- Extended Berkeley Packet Filters
  - But the tracing infrastructure can also exploit it



# eBPF

## Why eBPF?

- Minimal overhead
- Programmability
  - Compute customized event statistics
  - Perform in-kernel data aggregation
- Production-safe
  - Checks for unsafe code before execution
  - Execution in a secure VM
- Maps
  - Maintain state across events
  - Exchange data with Userspace

# BCC

## What is BCC?

- eBPF Compiler Collection
  - Toolkit for creating eBPF-based tracing scripts
- Provides a Python API
  - Simplifies tasks
    - Event creation
    - Compilation of eBPF programs
    - Loading and attaching eBPF programs
    - Access to maps
- Set of useful, readily-available tools
  - Trace a variety of kernel subsystems
  - Trace typical enterprise applications

# BCC

## Demos

- Attaching probes to a function
  - Print arguments
  - Print return value
- Filtering data by process
  - PID filter to capture data for a specific process
- Filtering data by setting thresholds
  - Find approximate execution time of a function
  - Find out when a function runs slower than a given threshold

<https://github.com/sandip4n/devconf.in-2018>

# Contact Us

Ravi Bangoria  
ravi.bangoria@linux.ibm.com

Sandipan Das  
sandipan@linux.ibm.com

# Thank You !

# References -- Perf

- Perf tool man pages
- `perf_event_open()` man page
- *perf Examples* by Brendan Gregg  
<http://www.brendangregg.com/perf.html>
- Perf wiki  
<https://perf.wiki.kernel.org/>
- *The Unofficial Linux Perf Events Web-Page* by Vince Weaver  
[http://web.eece.maine.edu/~vweaver/projects/perf\\_events/](http://web.eece.maine.edu/~vweaver/projects/perf_events/)

# References -- BCC and eBPF

- *A thorough introduction to eBPF* by Matt Fleming  
<https://lwn.net/Articles/740157/>
- *BPF - in-kernel virtual machine* by Alexei Starovoitov  
[https://events.linuxfoundation.org/sites/events/files/slides/bpf\\_collabsummit\\_2015feb20.pdf](https://events.linuxfoundation.org/sites/events/files/slides/bpf_collabsummit_2015feb20.pdf)
- *Linux BPF Superpowers* by Brendan Gregg  
<http://www.slideshare.net/brendangregg/linux-bpf-superpowers>
- *Linux Enhanced BPF (eBPF) Tracing Tools* by Brendan Gregg  
<http://www.brendangregg.com/ebpf.html>
- *BPF: tracing and more* by Brendan Gregg  
<http://www.slideshare.net/brendangregg/bpf-tracing-and-more>

# Legal Statement

- This work represents the view of the authors and does not necessarily represent the view of the employers (IBM Corporation).
- IBM and IBM (Logo) are trademarks or registered trademarks of International Business Machines in United States and/or other countries.
- Linux is a registered trademark of Linus Torvalds.
- Other company, product and service names may be trademarks or service marks of others.



# Linux Tracing Tools -- Backup

- Perf
- BCC (eBPF)
- Strace
- Ftrace
- Systemtap
- LTTng
- PCP
- Gprof
- Oprofile